

```
import numpy as np # linear algebra
import pandas as pd
```

Objective

The objective of this project is to build and evaluate a Convolutional Neural Network (CNN) model to classify handwritten digits using the MNIST dataset. The aim is to achieve high accuracy in recognizing digits from 0 to 9, thereby demonstrating the effectiveness of CNNs in image recognition tasks.

This project involves several key steps:

Data Loading and Preprocessing:

Load the MNIST dataset, which contains 60,000 training images and 10,000 test images of handwritten digits. Normalize the pixel values to range between 0 and 1. Exploratory Data Analysis:

Visualize random samples from the dataset to understand the input data. Model Building:

Construct a Sequential CNN model with two convolutional layers, each followed by a max-pooling layer. Add a fully connected (dense) layer and an output layer with softmax activation to classify the digits. Model Training:

Compile the model using the Adam optimizer and categorical cross-entropy loss function. Train the model on the training data for 10 epochs, using a batch size of 64, and validate it on the test data. Model Evaluation:

Evaluate the model's performance on the test dataset to determine its accuracy and loss. Visualize the training and validation accuracy and loss over epochs to understand the model's learning curve.

✓ Data Loading and Preprocessing:

```
import tensorflow as tf
from keras.datasets import mnist
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout
from tensorflow.keras.utils import to_categorical
import matplotlib.pyplot as plt
import numpy as np
import os
for dirname, _, filenames in os.walk('/content/Mnist dataset.zip'):
    for filename in filenames:
        print(os.path.join(dirname, filename))
```


```
(train_images, train_labels), (test_images, test_labels) = mnist.load_data()
train_images, test_images = train_images / 255.0, test_images / 255.0
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>
11490434/11490434 — 0s 0us/step

✓ Exploratory Data Analysis:

```
random_index = np.random.randint(0, len(train_images))  
random_image = train_images[random_index]  
random_label = train_labels[random_index]
```

```
plt.imshow(random_image)  
plt.axis('off')  
plt.show()
```



```
random_image.shape
```

(28, 28)

```
labels = str(list(range(1,11)))
```


```
# Selecting 12 random images  
num_samples = 12  
random_indices = np.random.choice(train_images.shape[0], num_samples, replace=False)  
sample_images = train_images[random_indices]  
sample_labels = train_labels[random_indices]
```

sample_labels

```
array([6, 3, 5, 9, 6, 9, 1, 8, 0, 6, 6, 5], dtype=uint8)
```

```
# Plotting
plt.figure(figsize=(10, 10))
for i in range(num_samples):
    plt.subplot(4, 4, i + 1)
    plt.imshow(sample_images[i])
    # plt.title(i)

    plt.title(int(sample_labels[i])) # Ensure sample_labels[i][0] is an index within the range of labels
    plt.axis('off')
plt.tight_layout()
plt.show()
```



```
train_labels = to_categorical(train_labels)
test_labels = to_categorical(test_labels)
```


```
train_images.shape
```

(60000, 28, 28)

```
# Build Convolutional Neural Network Model
model = tf.keras.models.Sequential([
    # Ist Convolutional Layer
    tf.keras.layers.Conv2D(32, (3, 3), strides=(1, 1), activation='relu', input_shape=(28, 28, 1)),
    tf.keras.layers.MaxPooling2D((2, 2)),












    # 2nd convolutional layer
    tf.keras.layers.Conv2D(64, (3, 3), strides=(1, 1), activation='relu'),
    tf.keras.layers.MaxPooling2D((2, 2)),

    # ANN
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(32, activation='relu'),
    tf.keras.layers.Dense(10, activation='softmax')
])
```

 /usr/local/lib/python3.10/dist-packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When super().__init__(activity_regularizer=activity_regularizer, **kwargs)

```
# Modeling
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model.fit(train_images, train_labels, epochs=10, batch_size=64,
                    validation_data=(test_images, test_labels))
```

 Epoch 1/10
938/938  **73s** 74ms/step - accuracy: 0.8670 - loss: 0.4507 - val_accuracy: 0.9760 - val_loss: 0.0718
 Epoch 2/10
938/938  **63s** 67ms/step - accuracy: 0.9800 - loss: 0.0647 - val_accuracy: 0.9820 - val_loss: 0.0568
 Epoch 3/10
938/938  **76s** 60ms/step - accuracy: 0.9861 - loss: 0.0447 - val_accuracy: 0.9889 - val_loss: 0.0348
 Epoch 4/10
938/938  **83s** 62ms/step - accuracy: 0.9895 - loss: 0.0345 - val_accuracy: 0.9886 - val_loss: 0.0322
 Epoch 5/10
938/938  **80s** 60ms/step - accuracy: 0.9925 - loss: 0.0258 - val_accuracy: 0.9883 - val_loss: 0.0374
 Epoch 6/10
938/938  **56s** 60ms/step - accuracy: 0.9926 - loss: 0.0232 - val_accuracy: 0.9898 - val_loss: 0.0299
 Epoch 7/10
938/938  **65s** 69ms/step - accuracy: 0.9940 - loss: 0.0176 - val_accuracy: 0.9896 - val_loss: 0.0317
 Epoch 8/10
938/938  **73s** 60ms/step - accuracy: 0.9961 - loss: 0.0126 - val_accuracy: 0.9906 - val_loss: 0.0314
 Epoch 9/10
938/938  **57s** 61ms/step - accuracy: 0.9960 - loss: 0.0118 - val_accuracy: 0.9909 - val_loss: 0.0311
 Epoch 10/10
938/938  **79s** 58ms/step - accuracy: 0.9968 - loss: 0.0111 - val_accuracy: 0.9907 - val_loss: 0.0338

✓ Displaying the model architecture and the number of parameters.


```
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_1 (Conv2D)	(None, 11, 11, 64)	18,496
max_pooling2d_1 (MaxPooling2D)	(None, 5, 5, 64)	0
flatten (Flatten)	(None, 1600)	0
dense (Dense)	(None, 32)	51,232
dense_1 (Dense)	(None, 10)	330

Total params: 211,136 (824.75 KB)
Trainable params: 70,378 (274.91 KB)
Non-trainable params: 0 (0.00 B)
Optimizer params: 140,758 (549.84 KB)

Model Building:



```

import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Flatten
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.losses import SparseCategoricalCrossentropy

# Load and preprocess the dataset (using MNIST as an example)
(train_images, train_labels), (test_images, test_labels) = tf.keras.datasets.mnist.load_data()

# Normalize the images to the range [0, 1]
train_images = train_images / 255.0
test_images = test_images / 255.0


# Define the model
model = Sequential([
    Flatten(input_shape=(28, 28)), # Flatten the input data
    Dense(64, activation='relu'),
    Dense(64, activation='relu'),
    Dense(10, activation='softmax')
])


# Compile the model
model.compile(optimizer=Adam(),
              loss=SparseCategoricalCrossentropy(),
              metrics=['accuracy'])

# Evaluate the model
loss, accuracy = model.evaluate(train_images, train_labels)

# Print the results
print(f"The model accuracy is: {accuracy}\nThe model loss is: {loss}")

```

 /usr/local/lib/python3.10/dist-packages/keras/src/layers/reshaping/flatten.py:37: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using `super().__init__(**kwargs)`

1875/1875  **3s** 2ms/step - accuracy: 0.1047 - loss: 2.3639

The model accuracy is: 0.10573333501815796



The model loss is: 2.36128830909729

✓ Model Evaluation:

```

loss, accuracy = model.evaluate(train_images, train_labels)
print(f"The model accuracy is : {accuracy} \n the model loss : {loss}")

```

 **1875/1875**  **3s** 2ms/step - accuracy: 0.1047 - loss: 2.3639

The model accuracy is : 0.10573333501815796

the model loss : 2.36128830909729

```
import matplotlib.pyplot as plt
```


```
history = model.fit(train_images, train_labels, epochs=10, validation_split=0.2)
```

```
Epoch 1/10
1500/1500 ————— 7s 3ms/step - accuracy: 0.8520 - loss: 0.5176 - val_accuracy: 0.9475 - val_loss: 0.1757
Epoch 2/10
1500/1500 ————— 5s 3ms/step - accuracy: 0.9541 - loss: 0.1483 - val_accuracy: 0.9587 - val_loss: 0.1368
Epoch 3/10
1500/1500 ————— 6s 4ms/step - accuracy: 0.9667 - loss: 0.1059 - val_accuracy: 0.9643 - val_loss: 0.1264
Epoch 4/10
1500/1500 ————— 11s 4ms/step - accuracy: 0.9750 - loss: 0.0806 - val_accuracy: 0.9669 - val_loss: 0.1137
Epoch 5/10
1500/1500 ————— 6s 4ms/step - accuracy: 0.9789 - loss: 0.0671 - val_accuracy: 0.9678 - val_loss: 0.1134
Epoch 6/10
1500/1500 ————— 6s 4ms/step - accuracy: 0.9836 - loss: 0.0526 - val_accuracy: 0.9712 - val_loss: 0.1017
Epoch 7/10
1500/1500 ————— 9s 3ms/step - accuracy: 0.9864 - loss: 0.0417 - val_accuracy: 0.9696 - val_loss: 0.1077
Epoch 8/10
1500/1500 ————— 9s 3ms/step - accuracy: 0.9876 - loss: 0.0382 - val_accuracy: 0.9688 - val_loss: 0.1139
Epoch 9/10
1500/1500 ————— 5s 4ms/step - accuracy: 0.9895 - loss: 0.0344 - val_accuracy: 0.9719 - val_loss: 0.1064
Epoch 10/10
1500/1500 ————— 9s 3ms/step - accuracy: 0.9909 - loss: 0.0283 - val_accuracy: 0.9708 - val_loss: 0.1123
```

✓ Visualization of Training History:

```
# Grapgh
```

```
plt.plot(history.history['accuracy'],label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label = 'Validation Accuracy')
plt.xlabel("Epochs")
plt.ylabel('Accuracy')
plt.legend()
plt.show()
```

```
plt.plot(history.history['loss'],label='Training Loss')
plt.plot(history.history['val_loss'],label='Validation Loss')
plt.xlabel("Epoche")
plt.ylabel("Loss")
plt.legend()
```